

Machine Learning Foundations

(機器學習基石)



Lecture 3: Types of Learning

Hsuan-Tien Lin (林軒田)

htlin@csie.ntu.edu.tw

Department of Computer Science
& Information Engineering

National Taiwan University
(國立台灣大學資訊工程系)



Roadmap

① When Can Machines Learn?

Lecture 2: Learning to Answer Yes/No

PLA \mathcal{A} takes **linear separable** \mathcal{D} and **perceptrons** \mathcal{H} to get **hypothesis** g

Lecture 3: Types of Learning

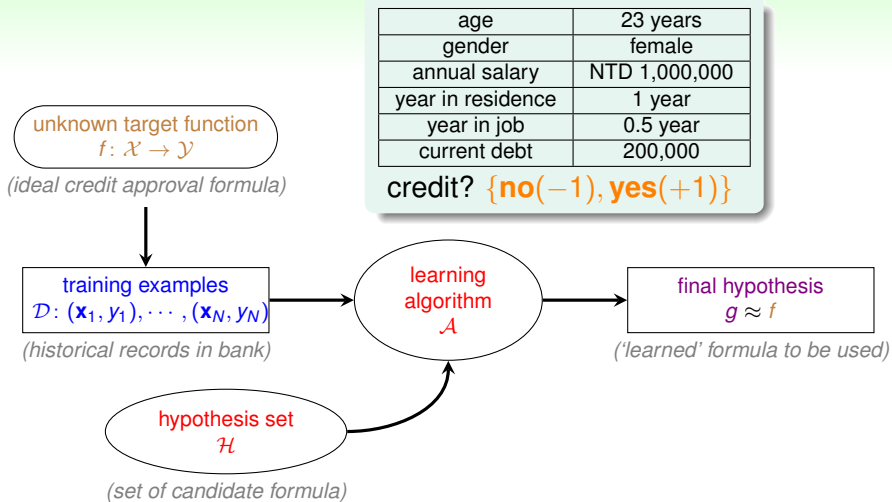
- Learning with Different Output Space \mathcal{Y}
- Learning with Different Data Label y_n
- Learning with Different Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$
- Learning with Different Input Space \mathcal{X}

② Why Can Machines Learn?

③ How Can Machines Learn?




④ How Can Machines Learn Better?

Credit Approval Problem Revisited



$\mathcal{Y} = \{-1, +1\}$: **binary classification**


More Binary Classification Problems



- credit **approve/disapprove**
- email **spam/non-spam**
- patient **sick/not sick**
- ad **profitable/not profitable**
- answer **correct/incorrect** (KDDCup 2010)

core and important problem with
many tools as **building block of other tools**

Multiclass Classification: Coin Recognition Problem



- classify US coins (1c, 5c, 10c, 25c) by (size, mass)
- $\mathcal{Y} = \{1c, 5c, 10c, 25c\}$, or
 $\mathcal{Y} = \{1, 2, \dots, K\}$ (**abstractly**)
- binary classification: special case with $K = 2$

Other Multiclass Classification Problems

- written digits $\Rightarrow 0, 1, \dots, 9$
- pictures \Rightarrow apple, orange, strawberry
- emails \Rightarrow spam, primary, social, promotion, update (Google)

many applications in practice,
especially for 'recognition'

Regression: Patient Recovery Prediction Problem

- binary classification: patient features \Rightarrow sick or not
- multiclass classification: patient features \Rightarrow which type of cancer
- regression: patient features \Rightarrow **how many days before recovery**
- $\mathcal{Y} = \mathbb{R}$ or $\mathcal{Y} = [\text{lower}, \text{upper}] \subset \mathbb{R}$ (bounded regression)
— **deeply studied in statistics**

Other Regression Problems

- company data \Rightarrow stock price
- climate data \Rightarrow temperature

also core and important with many ‘statistical’
tools as **building block of other tools**

Structured Learning: Sequence Tagging Problem

除了分类、回归问题，还有结构化学习：序列标注问题。

$\underbrace{I}_{\text{pronoun}} \underbrace{\text{love}}_{\text{verb}} \underbrace{ML}_{\text{noun}}$

- multiclass classification: word \Rightarrow word class
- structured learning:
sentence \Rightarrow structure (class of each word)
- $\mathcal{Y} = \{PVN, PVP, NVN, PV, \dots\}$, not including VVVVV
- huge multiclass classification problem
(structure \equiv hyperclass) **without 'explicit' class definition**

Other Structured Learning Problems

- protein data \Rightarrow protein folding
- speech data \Rightarrow speech parse tree


a fancy but complicated learning problem

Mini Summary

根据输出 y 的不同，可
以大致分为以下几类问
题：

Learning with Different Output Space \mathcal{Y}

- **binary classification**: $\mathcal{Y} = \{-1, +1\}$
- **multiclass classification**: $\mathcal{Y} = \{1, 2, \dots, K\}$
- **regression**: $\mathcal{Y} = \mathbb{R}$
- **structured learning**: $\mathcal{Y} = \text{structures}$
- ... and a lot more!!



core tools: binary classification and regression

Fun Time

What is this learning problem?

The entrance system of the school gym, which does automatic face recognition based on machine learning, is built to charge four different groups of users differently: Staff, Student, Professor, Other. What type of learning problem best fits the need of the system?

- ① binary classification
- ② multiclass classification
- ③ regression
- ④ structured learning

Fun Time

What is this learning problem?

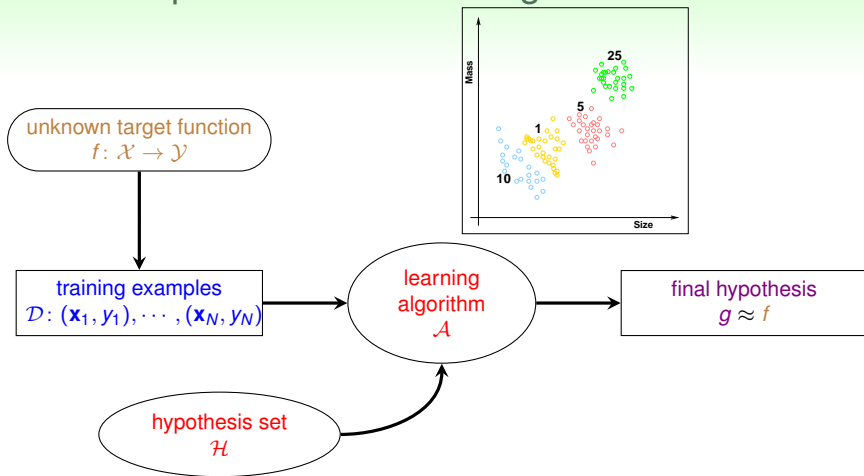
The entrance system of the school gym, which does automatic face recognition based on machine learning, is built to charge four different groups of users differently: Staff, Student, Professor, Other. What type of learning problem best fits the need of the system?

- ① binary classification
- ② multiclass classification
- ③ regression
- ④ structured learning

Reference Answer: ②


There is an 'explicit' \mathcal{Y} that contains four classes.

Supervised: Coin Recognition Revisited




supervised learning:
every \mathbf{x}_n comes with corresponding y_n

Unsupervised: Coin Recognition without y_n



supervised multiclass classification




unsupervised multiclass classification
 \iff 'clustering'

Other Clustering Problems


- articles \Rightarrow topics
- consumer profiles \Rightarrow consumer groups

clustering: a challenging but useful problem

Unsupervised: Coin Recognition without y_n



supervised multiclass classification



unsupervised multiclass classification
 \iff 'clustering'

Other Clustering Problems

- articles \Rightarrow topics
- consumer profiles \Rightarrow consumer groups

clustering: a challenging but useful problem

Unsupervised: Learning without y_n


Other Unsupervised Learning Problems

- clustering: $\{\mathbf{x}_n\} \Rightarrow \text{cluster}(\mathbf{x})$
(\approx 'unsupervised multiclass classification')
—i.e. articles \Rightarrow topics
- **density estimation**: $\{\mathbf{x}_n\} \Rightarrow \text{density}(\mathbf{x})$
(\approx 'unsupervised bounded regression')
—i.e. traffic reports with location \Rightarrow dangerous areas
- **outlier detection**: $\{\mathbf{x}_n\} \Rightarrow \text{unusual}(\mathbf{x})$
(\approx extreme 'unsupervised binary classification')
—i.e. Internet logs \Rightarrow intrusion alert
- ... and a lot more!!


密度估计和异常检测属于无监督学习的内容。

unsupervised learning: diverse, with possibly very different performance goals


Semi-supervised: Coin Recognition with Some y_n



supervised



semi-supervised



unsupervised (clustering)

Other Semi-supervised Learning Problems

- face images with a few labeled \Rightarrow face identifier (Facebook)
- medicine data with a few labeled \Rightarrow medicine effect predictor

semi-supervised learning: leverage unlabeled data to avoid 'expensive' labeling

Reinforcement Learning


a 'very different' but natural way of learning

Teach Your Dog: Say 'Sit Down'

The dog **pees** on the ground. 尿

BAD DOG. THAT'S A VERY WRONG ACTION.

- cannot easily show the dog that $y_n = \text{sit}$ when $\mathbf{x}_n = \text{'sit down'}$
- but can 'punish' to say $\tilde{y}_n = \text{pee is wrong}$



Other Reinforcement Learning Problems Using $(\mathbf{x}, \tilde{y}, \text{goodness})$

- (customer, ad choice, ad click earning) \Rightarrow ad system
- (cards, strategy, winning amount) \Rightarrow black jack agent

reinforcement: learn with '**partial/implicit information**' (often sequentially)

Reinforcement Learning


a 'very different' but natural way of learning

Teach Your Dog: Say 'Sit Down'

The dog sits down.

Good Dog. Let me give you some cookies.

- still cannot show $y_n = \text{sit}$ when $\mathbf{x}_n = \text{'sit down'}$
- but can 'reward' to say $\tilde{y}_n = \text{sit}$ is good



Other Reinforcement Learning Problems Using $(\mathbf{x}, \tilde{y}, \text{goodness})$

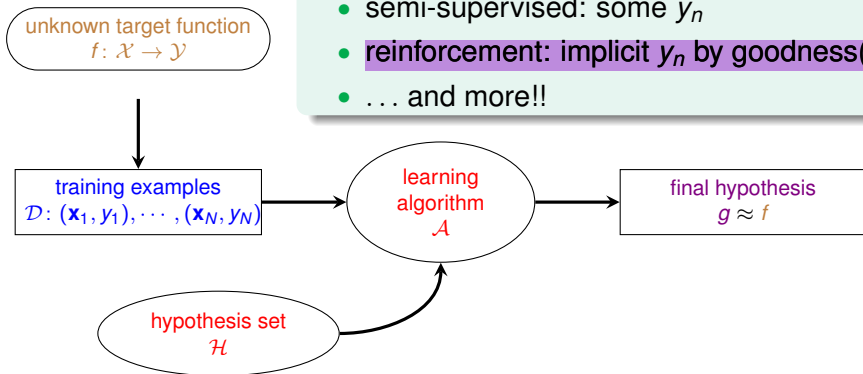
- (customer, ad choice, ad click earning) \Rightarrow ad system
- (cards, strategy, winning amount) \Rightarrow black jack agent

reinforcement: learn with '**partial/implicit information**' (often sequentially)

Mini Summary

Learning with Different Data Label y_n

- **supervised**: all y_n
- unsupervised: no y_n
- semi-supervised: some y_n
- reinforcement: implicit y_n by goodness(\tilde{y}_n)
- ... and more!!



core tool: supervised learning

Fun Time

What is this learning problem?

To build a tree recognition system, a company decides to gather one million of pictures on the Internet. Then, it asks each of the 10 company members to view 100 pictures and record whether each picture contains a tree. The pictures and records are then fed to a learning algorithm to build the system. What type of learning problem does the algorithm need to solve?

- ① supervised
- ② unsupervised
- ③ semi-supervised
- ④ reinforcement

Fun Time

What is this learning problem?

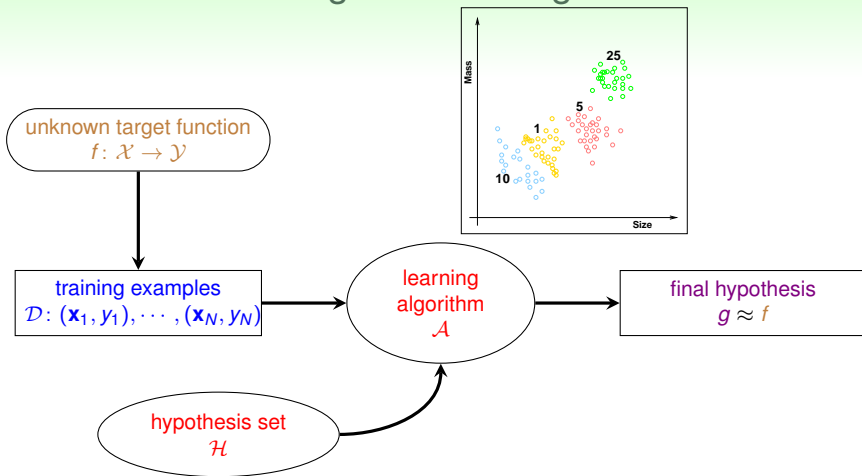
To build a tree recognition system, a company decides to gather one million of pictures on the Internet. Then, it asks each of the 10 company members to view 100 pictures and record whether each picture contains a tree. The pictures and records are then fed to a learning algorithm to build the system. What type of learning problem does the algorithm need to solve?

- ① supervised
- ② unsupervised
- ③ semi-supervised
- ④ reinforcement

Reference Answer: ③



The 1,000 records are the labeled (\mathbf{x}_n, y_n) ; the other 999,000 pictures are the unlabeled \mathbf{x}_n .

Batch Learning: Coin Recognition Revisited



batch supervised multiclass classification:
learn from **all known** data

More Batch Learning Problems



- batch of (email, spam?) \Rightarrow spam filter
- batch of (patient, cancer) \Rightarrow cancer classifier
- batch of patient data \Rightarrow group of patients

batch learning: **a very common protocol**

Online: Spam Filter that 'Improves'

- batch spam filter:
learn with known (email, spam?) pairs, and predict with fixed g
- **online** spam filter, which **sequentially**:
 - 1 observe an email \mathbf{x}_t
 - 2 predict spam status with current $g_t(\mathbf{x}_t)$
 - 3 receive 'desired label' y_t from user, and then update g_t with (\mathbf{x}_t, y_t)

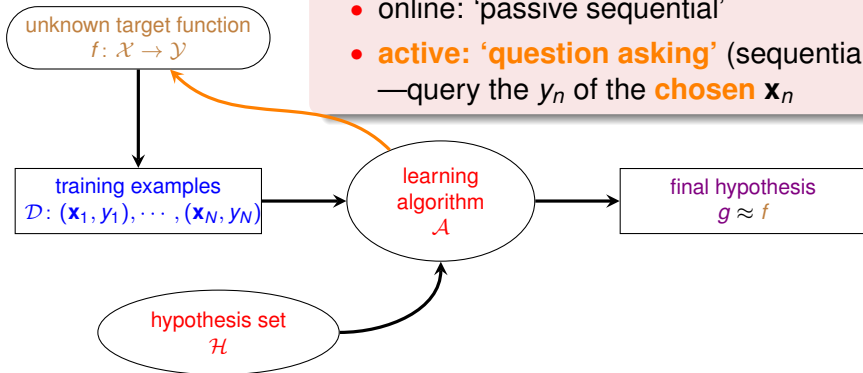
Connection to What We Have Learned

- PLA can be easily adapted to online protocol (how?)
- reinforcement learning is often done online (why?)

online: hypothesis 'improves' through receiving
data instances **sequentially**

主动学习 **Active Learning: Learning by 'Asking'****Protocol \Leftrightarrow Learning Philosophy**

- batch: 'duck feeding'
- online: 'passive sequential'
- **active: 'question asking'** (sequentially)
—query the y_n of the **chosen** \mathbf{x}_n

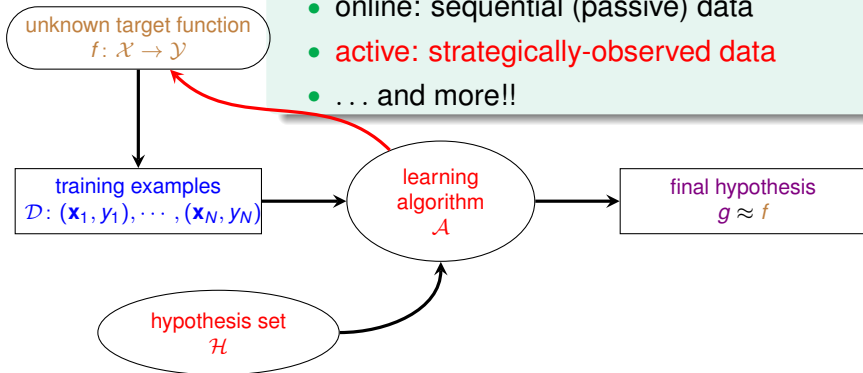


**active: improve hypothesis with fewer labels
(hopefully) by asking questions strategically**

Mini Summary

Learning with Different Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$

- **batch**: all known data
- online: sequential (passive) data
- **active**: strategically-observed data
- ... and more!!



core protocol: batch

Fun Time

What is this learning problem?

A photographer has 100,000 pictures, each containing one baseball player. He wants to automatically categorize the pictures by its player inside. He starts by categorizing 1,000 pictures by himself, and then writes an algorithm that tries to categorize the other pictures if it is 'confident' on the category while pausing for (& learning from) human input if not. What protocol best describes the nature of the algorithm?

- ① batch
- ② online
- ③ active
- ④ random

Fun Time

What is this learning problem?

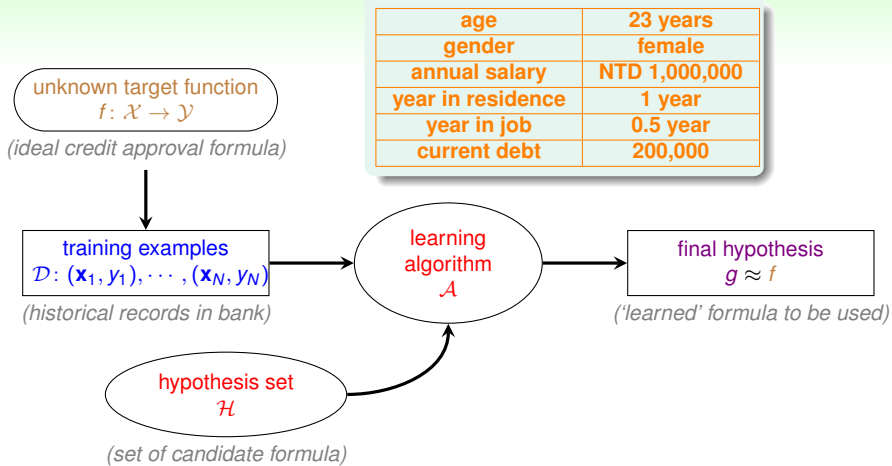
A photographer has 100,000 pictures, each containing one baseball player. He wants to automatically categorize the pictures by its player inside. He starts by categorizing 1,000 pictures by himself, and then writes an algorithm that tries to categorize the other pictures if it is 'confident' on the category while pausing for (& learning from) human input if not. What protocol best describes the nature of the algorithm?

- ① batch
- ② online
- ③ active
- ④ random

Reference Answer: ③

The algorithm takes a active but naïve strategy: ask when 'confused'. **You should probably do the same when taking a class. :-)**


Credit Approval Problem Revisited



concrete features: each dimension of $\mathcal{X} \subseteq \mathbb{R}^d$ represents ‘sophisticated physical meaning’


More on Concrete Features

- **(size, mass)** for coin classification
- **customer info** for credit approval
- **patient info** for cancer diagnosis
- often including 'human intelligence' on the learning task



concrete features: the 'easy' ones for ML


Raw Features: Digit Recognition Problem (1/2)



- digit recognition problem: features \Rightarrow meaning of digit
- a typical supervised multiclass classification problem

Raw Features: Digit Recognition Problem (2/2)

by Concrete Features



$\mathbf{x} = (\text{symmetry, density})$

by Raw Features

- 16 by 16 gray image $\mathbf{x} \equiv (0, 0, 0.9, 0.6, \dots) \in \mathbb{R}^{256}$
- ‘**simple** physical meaning’; thus more difficult for ML than concrete features

Other Problems with Raw Features

- image pixels, speech signal, etc.

raw features: often need human or machines
to **convert to concrete ones**

Abstract Features: Rating Prediction Problem

Rating Prediction Problem (KDDCup 2011)

- given previous (userid, itemid, rating) tuples, predict the rating that some userid would give to itemid?
- a regression problem with $\mathcal{Y} \subseteq \mathbb{R}$ as rating and $\mathcal{X} \subseteq \mathbb{N} \times \mathbb{N}$ as (userid, itemid)
- ‘no physical meaning’; thus even more difficult for ML

Other Problems with Abstract Features

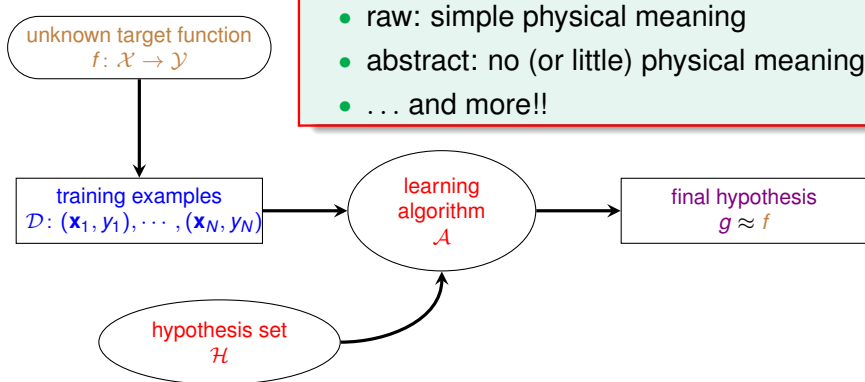
- student ID in online tutoring system (KDDCup 2010)
- advertisement ID in online ad system

abstract: again need ‘**feature conversion**/extraction/construction’

Mini Summary

Learning with Different Input Space \mathcal{X}

- **concrete**: sophisticated (and related) physical meaning
- raw: simple physical meaning
- abstract: no (or little) physical meaning
- ... and more!!



'easy' input: concrete

Fun Time

What features can be used?

Consider a problem of building an online image advertisement system that shows the users the most relevant images. What features can you choose to use?

- 1 concrete
- 2 concrete, raw
- 3 concrete, abstract
- 4 concrete, raw, abstract

Fun Time

What features can be used?

Consider a problem of building an online image advertisement system that shows the users the most relevant images. What features can you choose to use?

- ① concrete
- ② concrete, raw
- ③ concrete, abstract
- ④ concrete, raw, abstract

Reference Answer: ④

concrete user features, raw image features,
and maybe abstract user/image IDs

Summary

1 When Can Machines Learn?

Lecture 2: Learning to Answer Yes/No

Lecture 3: Types of Learning

- Learning with Different Output Space \mathcal{Y}
[classification], [regression], structured
 - Learning with Different Data Label y_n
[supervised], un/semi-supervised, reinforcement
 - Learning with Different Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$
[batch], online, active
 - Learning with Different Input Space \mathcal{X}
[concrete], raw, abstract
- **next: learning is impossible?!**
- 2 Why Can Machines Learn?
 - 3 How Can Machines Learn?
 - 4 How Can Machines Learn Better?