


# Machine Learning Foundations

## (機器學習基石)



### Lecture 13: Hazard of Overfitting

Hsuan-Tien Lin (林軒田)

htlin@csie.ntu.edu.tw

Department of Computer Science  
& Information Engineering

National Taiwan University  
(國立台灣大學資訊工程系)



# Roadmap

- 1 When Can Machines Learn?
- 2 Why Can Machines Learn?
- 3 How Can Machines Learn?

## Lecture 12: Nonlinear Transform

**nonlinear**  $\square$  via **nonlinear feature transform  $\phi$**   
plus **linear**  $\square$  with price of **model complexity**


- 4 How Can Machines Learn **Better**?

## Lecture 13: Hazard of Overfitting

- What is Overfitting?
- The Role of Noise and Data Size
- Deterministic Noise
- Dealing with Overfitting

# Bad Generalization

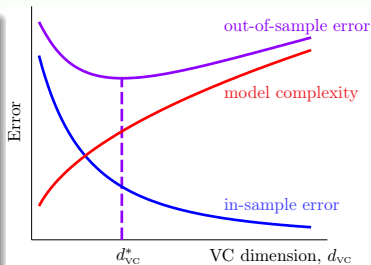
- regression for  $x \in \mathbb{R}$  with  $N = 5$  examples
- target  $f(x)$  = 2nd order polynomial
- label  $y_n = f(x_n) + \text{very small noise}$
- linear regression in  $\mathcal{Z}$ -space +  $\Phi$  = 4th order polynomial
- unique solution passing all examples  $\implies E_{\text{in}}(g) = 0$
- $E_{\text{out}}(g)$  huge



bad generalization: low  $E_{\text{in}}$ , high  $E_{\text{out}}$

# Bad Generalization and Overfitting

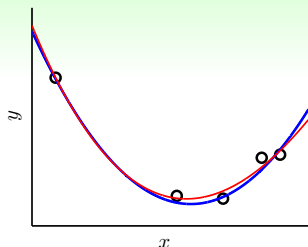
- take  $d_{VC} = 1126$  for learning:  
bad generalization  
—( $E_{out} - E_{in}$ ) large
- switch from  $d_{VC} = d_{VC}^*$  to  $d_{VC} = 1126$ :  
**overfitting**  
— $E_{in} \downarrow$ ,  $E_{out} \uparrow$
- switch from  $d_{VC} = d_{VC}^*$  to  $d_{VC} = 1$ :  
**underfitting**  
— $E_{in} \uparrow$ ,  $E_{out} \uparrow$




bad generalization: low  $E_{in}$ , high  $E_{out}$ ;  
**overfitting**: lower  $E_{in}$ , higher  $E_{out}$

两者的不同：bad generalization是指一个点的状态，是结果  
 Overfitting是描述一个过程， $E_{in}$ 变低， $E_{out}$ 变高。

# Cause of Overfitting: A Driving Analogy



'good fit'



overfit

learning

overfit

use excessive  $d_{VC}$

**noise**

**limited data size  $N$**

driving

commit a car accident

'drive too fast'

**bumpy** road 凹凸不平的

limited observations about road condition

next: how does **noise** & **data size** affect overfitting?

## Fun Time

Based on our discussion, for data of fixed size, which of the following situation is relatively of the lowest risk of overfitting?

- ① small noise, fitting from small  $d_{VC}$  to median  $d_{VC}$
- ② small noise, fitting from small  $d_{VC}$  to large  $d_{VC}$
- ③ large noise, fitting from small  $d_{VC}$  to median  $d_{VC}$
- ④ large noise, fitting from small  $d_{VC}$  to large  $d_{VC}$

## Fun Time

Based on our discussion, for data of fixed size, which of the following situation is relatively of the lowest risk of overfitting?


- ① small noise, fitting from small  $d_{VC}$  to median  $d_{VC}$
- ② small noise, fitting from small  $d_{VC}$  to large  $d_{VC}$
- ③ large noise, fitting from small  $d_{VC}$  to median  $d_{VC}$
- ④ large noise, fitting from small  $d_{VC}$  to large  $d_{VC}$

Reference Answer: ①


Two causes of overfitting are noise and excessive  $d_{VC}$ . So if both are relatively 'under control', the risk of overfitting is smaller.

## Case Study (1/2)

10-th order target function  
+ noise



50-th order target function  
noiselessly



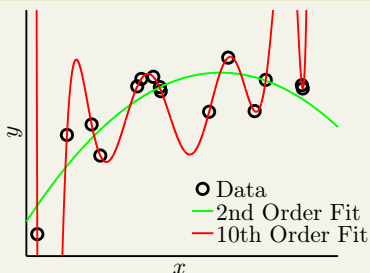
overfitting from best  $g_2 \in \mathcal{H}_2$  to best  $g_{10} \in \mathcal{H}_{10}$ ?



## Case Study (2/2)

10-th order target function  
+ noise

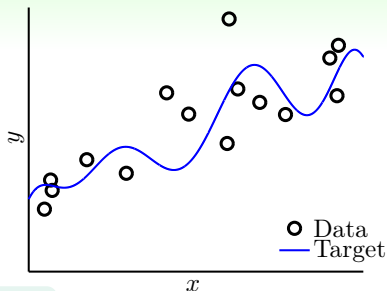
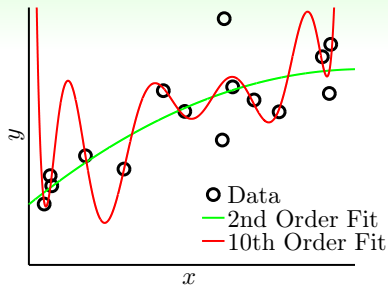
	$g_2 \in \mathcal{H}_2$	$g_{10} \in \mathcal{H}_{10}$
$E_{\text{in}}$	0.050	0.034
$E_{\text{out}}$	0.127	9.00

50-th order target function  
noiselessly

	$g_2 \in \mathcal{H}_2$	$g_{10} \in \mathcal{H}_{10}$
$E_{\text{in}}$	0.029	0.00001
$E_{\text{out}}$	0.120	7680

overfitting from  $g_2$  to  $g_{10}$ ? **both yes!**

# Irony of Two Learners





- learner **Overfit**: pick  $g_{10} \in \mathcal{H}_{10}$
- learner **Restrict**: pick  $g_2 \in \mathcal{H}_2$
- when both **know that target = 10th**  
 —  $R$  'gives up' ability to fit

but  $R$  **wins in  $E_{out}$**  a lot!

philosophy: **concession** for **advantage**? :-) 以退为进

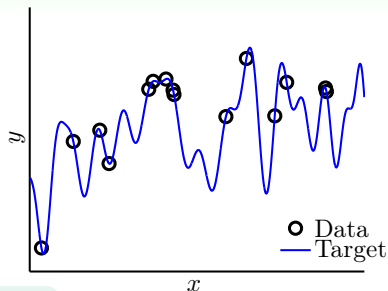
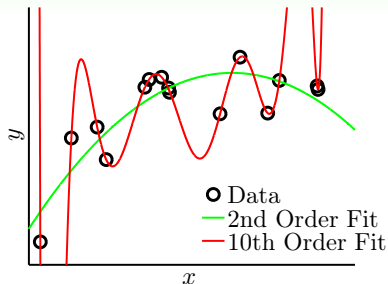
## Learning Curves Revisited



- $\mathcal{H}_{10}$ : lower  $\overline{E_{out}}$  when  $N \rightarrow \infty$ ,  
but much larger generalization error for small  $N$
- gray area:  $O$  overfits! ( $\overline{E_{in}} \downarrow$ ,  $\overline{E_{out}} \uparrow$ )

在样本数量不够多时，反而简单的模型效果要更好  
 $R$  always **wins in**  $E_{out}$  if  $N$  small!

# The 'No Noise' Case



- learner **Overfit**: pick  $g_{10} \in \mathcal{H}_{10}$
- learner **Restrict**: pick  $g_2 \in \mathcal{H}_2$
- when both **know that there is no noise** —  $R$  still wins

is there really **no noise**?  
'target complexity' acts like noise

# Fun Time

When having limited data, in which of the following case would learner  $R$  perform better than learner  $O$ ?

- ① limited data from a 10-th order target function with some noise
- ② limited data from a 1126-th order target function with no noise
- ③ limited data from a 1126-th order target function with some noise
- ④ all of the above

## Fun Time

When having limited data, in which of the following case would learner  $R$  perform better than learner  $O$ ?

- ① limited data from a 10-th order target function with some noise
- ② limited data from a 1126-th order target function with no noise
- ③ limited data from a 1126-th order target function with some noise
- ④ all of the above

Reference Answer: ④

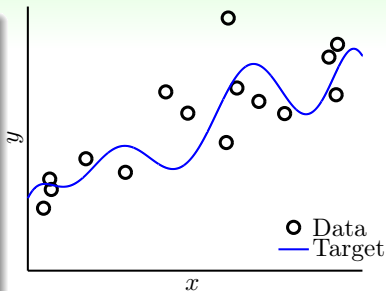
We discussed about ① and ②, but you shall be able to '**generalize**' :-)) that  $R$  also wins in the more difficult case of ③.

# A Detailed Experiment

$$y = f(x) + \epsilon$$

$$\sim \text{Gaussian}\left(\underbrace{\sum_{q=0}^{Q_f} \alpha_q x^q}_{f(x)}, \sigma^2\right)$$

- Gaussian iid noise  $\epsilon$  with level  $\sigma^2$
- some 'uniform' distribution on  $f(x)$  with complexity level  $Q_f$
- data size  $N$




goal: **'overfit level'** for  
different  $(N, \sigma^2)$  and  $(N, Q_f)$ ?

# The Overfit Measure



- $g_2 \in \mathcal{H}_2$
- $g_{10} \in \mathcal{H}_{10}$
- $E_{\text{in}}(g_{10}) \leq E_{\text{in}}(g_2)$  for sure




**overfit measure**  $E_{\text{out}}(g_{10}) - E_{\text{out}}(g_2)$




# The Results


## impact of $\sigma^2$ versus $N$



## impact of $Q_f$ versus $N$

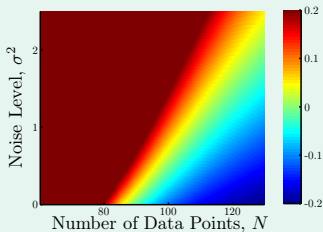


ring a bell? :-)

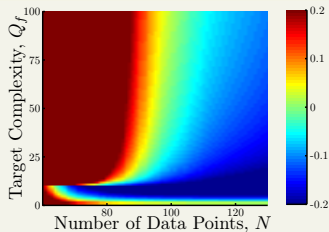


## Impact of Noise and Data Size

impact of  $\sigma^2$  versus  $N$ :  
**stochastic noise**



impact of  $Q_f$  versus  $N$ :  
**deterministic noise**



four reasons of serious overfitting:

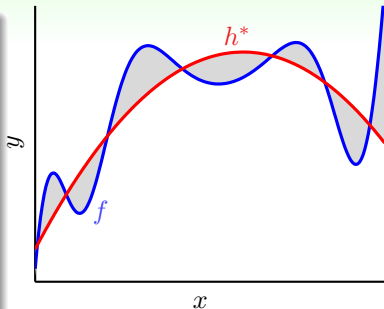
data size $N \downarrow$	overfit $\uparrow$
stochastic noise $\uparrow$	overfit $\uparrow$
deterministic noise $\uparrow$	overfit $\uparrow$
excessive power $\uparrow$	overfit $\uparrow$

**overfitting** 'easily' happens

如果目标函数太复杂的话，和noise没什么两样，我们把这种noise称为  
deterministic noise

## Deterministic Noise

- if  $f \notin \mathcal{H}$ : something of  $f$  cannot be captured by  $\mathcal{H}$
- **deterministic noise**: difference between best  $h^* \in \mathcal{H}$  and  $f$
- acts like 'stochastic noise'—not new to CS: **pseudo-random generator**
- difference to stochastic noise:
  - depends on  $\mathcal{H}$
  - fixed for a given  $\mathbf{x}$



philosophy: when teaching a kid,  
perhaps better not to use examples  
from a complicated target function? :-)

## Fun Time

Consider the target function being  $\sin(1126x)$  for  $x \in [0, 2\pi]$ . When  $x$  is uniformly sampled from the range, and we use all possible linear hypotheses  $h(x) = w \cdot x$  to approximate the target function with respect to the squared error, what is the level of deterministic noise for each  $x$ ?

- 1  $|\sin(1126x)|$
- 2  $|\sin(1126x) - x|$
- 3  $|\sin(1126x) + x|$
- 4  $|\sin(1126x) - 1126x|$

## Fun Time

Consider the target function being  $\sin(1126x)$  for  $x \in [0, 2\pi]$ . When  $x$  is uniformly sampled from the range, and we use all possible linear hypotheses  $h(x) = w \cdot x$  to approximate the target function with respect to the squared error, what is the level of deterministic noise for each  $x$ ?

- ①  $|\sin(1126x)|$
- ②  $|\sin(1126x) - x|$
- ③  $|\sin(1126x) + x|$
- ④  $|\sin(1126x) - 1126x|$

Reference Answer: ①


You can try a few different  $w$  and convince yourself that the best hypothesis  $h^*$  is  $h^*(x) = 0$ . The deterministic noise is the difference between  $f$  and  $h^*$ .

# Driving Analogy Revisited

learning	driving
overfit use excessive $d_{VC}$ noise limited data size $N$	commit a car accident 'drive too fast' bumpy road limited observations about road condition
<b>start from simple model</b> <b>data cleaning/pruning</b> <b>data hinting</b> <b>regularization</b> <b>validation</b>	drive slowly use more accurate road information exploit more road information put the brakes monitor the dashboard

all very **practical** techniques  
to combat overfitting

# Data Cleaning/Pruning




- if 'detect' the outlier **5** at the top by
  - too close to other  $\circ$ , or too far from other  $\times$
  - wrong by current classifier
  - ...
- possible action 1: correct the label (**data cleaning**) 数据修正
- possible action 2: remove the example (**data pruning**)

数据清理 (剪枝)

possibly helps, but **effect varies**

也许有用，但效果  
不尽相同

# Data Hinting



数据增强：可以使用平移，旋转等多种方式得到更多的数据

- slightly shifted/rotated digits carry the same meaning
- possible action: add **virtual examples** by shifting/rotating the given digits (**data hinting**)

possibly helps, but **watch out**  
 —**virtual example not**  $\stackrel{iid}{\sim} P(\mathbf{x}, y)$ !



# Fun Time

Assume we know that  $f(x)$  is symmetric for some 1D regression application. That is,  $f(x) = f(-x)$ . One possibility of using the knowledge is to consider symmetric hypotheses only. On the other hand, you can also generate virtual examples from the original data  $\{(x_n, y_n)\}$  as hints. What virtual examples suit your needs best?

- ①  $\{(x_n, -y_n)\}$
- ②  $\{(-x_n, -y_n)\}$
- ③  $\{(-x_n, y_n)\}$
- ④  $\{(2x_n, 2y_n)\}$

# Fun Time

Assume we know that  $f(x)$  is symmetric for some 1D regression application. That is,  $f(x) = f(-x)$ . One possibility of using the knowledge is to consider symmetric hypotheses only. On the other hand, you can also generate virtual examples from the original data  $\{(x_n, y_n)\}$  as hints. What virtual examples suit your needs best?

- ①  $\{(x_n, -y_n)\}$
- ②  $\{(-x_n, -y_n)\}$
- ③  $\{(-x_n, y_n)\}$
- ④  $\{(2x_n, 2y_n)\}$

Reference Answer: ③

We want the virtual examples to encode the invariance when  $x \rightarrow -x$ .

# Summary

- 1 When Can Machines Learn?
- 2 Why Can Machines Learn?
- 3 How Can Machines Learn?

## Lecture 12: Nonlinear Transform

- 4 How Can Machines Learn **Better**?

## Lecture 13: Hazard of Overfitting

- What is Overfitting?  
**lower  $E_{in}$  but higher  $E_{out}$**
- The Role of Noise and Data Size  
**overfitting 'easily' happens!**
- Deterministic Noise  
**what  $\mathcal{H}$  cannot capture acts like noise**
- Dealing with Overfitting  
**data cleaning/pruning/hinting, and more**

- **next: putting the brakes with regularization**